

Kevin Palm, Udacity Machine Learning Nanodegree - Project 4, 8/2016

**QUESTION 1:** Observe what you see with the agent's behavior as it takes random actions. Does the smartcab eventually make it to the destination? Are there any other interesting observations to note?

**ANSWER 1:** When the agent is taking random actions, I notice that the rewards are mostly negative. It looks like the agent runs a lot of red lights, cuts other drivers off, ect. While I was watching the agent never arrived at a target destination, but taking random actions could theoretically get the agent there. There would just be a lot of luck involved. I also notice that when the agent goes off the side of the map, it appears on the opposite side - so this grid is actually a flattened sphere.

**QUESTION 2:** What states have you identified that are appropriate for modeling the smartcab and environment? Why do you believe each of these states to be appropriate for this problem?

**ANSWER 2:** I've defined the state as a list: the light color, if there is any traffic surrounding the car, and which direction the planner is directing the agent.

My logic for including all of these components in state are:

- Light color - important for yielding
- oncoming front - also important for yielding
- oncoming left - yielding again
- oncoming right - yielding
- planner direction - important for getting to (or avoiding) the target

**OPTIONAL:** How many states in total exist for the smartcab in this environment? Does this number seem reasonable given that the goal of Q-Learning is to learn and make informed decisions about each state? Why or why not?

**OPTIONAL ANSWER:** There are 48 unique states possible given my definition of state, which seems reasonable for the problem and the amount of trials during which the agent gets to learn.

**QUESTION 3:** What changes do you notice in the agent's behavior when compared to the basic driving agent when random actions were always taken? Why is this behavior occurring?

**ANSWER 3:** The agent does start to seek out the target destinations. It also stops running red lights, and cutting off other drivers. This is occurring because that each action and reward are now being used to update the agent's policy concerning that state, so when the agent finds itself in the same or similar state again it can use the updated policies to pick the next action, rather than a random choice. The overall effect is that there's a lot less negative rewards getting incurred constantly, and the smartcab does start seeking out the destination.

**QUESTION 4:** Report the different values for the parameters tuned in your

basic implementation of Q-Learning. For which set of parameters does the agent perform best? How well does the final driving agent perform?

**ANSWER 4:** Here's a table of the parameters I tried and their performances on the last ten rounds (average of 10 runs).

Epsilon	Alpha	Gamma	Rewards per Turn	Destinations Reached
0.2	0.2	0.3	2.207367602	1
0.1	0.2	0.1	2.079277851	1
0.1	0.2	0.3	2.07014653	1
0.3	0.2	0.1	2.033722557	1
0.3	0.1	0.1	2.029828885	1
0.1	0.3	0.1	2.000196694	1
0.1	0.1	0.1	1.967406606	1
0.2	0.1	0.1	1.965554752	1
0.3	0.3	0.1	1.962549466	1
0.3	0.1	0.3	2.15948319	0.988888889
0.2	0.3	0.3	2.152544754	0.988888889
0.3	0.2	0.3	2.093428896	0.988888889
0.2	0.3	0.1	1.862837965	0.988888889
0.3	0.3	0.3	1.85189416	0.988888889
0.1	0.3	0.3	2.176400109	0.977777778
0.2	0.2	0.1	2.033498415	0.977777778
0.3	0.1	0.5	2.310587802	0.966666667
0.2	0.1	0.3	2.043346077	0.966666667
0.2	0.2	0.5	2.446405432	0.955555556
0.1	0.1	0.5	2.434614228	0.955555556
0.3	0.2	0.5	2.244029687	0.955555556
0.2	0.1	0.5	2.193300533	0.955555556
0.2	0.3	0.5	2.179456248	0.955555556
0.1	0.1	0.3	2.069022578	0.955555556
0.1	0.2	0.5	2.265488468	0.933333333
0.3	0.3	0.5	2.139192292	0.922222222
0.1	0.3	0.5	2.400442506	0.911111111

Using this table I selected the top combination which maximized destinations reached, and then maximized rewards per turn. The combination was epsilon=0.2, alpha=0.2, and gamma=0.3.

**QUESTION 5:** Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties? How would you describe an optimal policy for this problem?

**ANSWER 5:** Running the experiment 100 consecutive times and averaging together the trials showed that my selected tuning was closer to scores of "Average Rewards per Turn for the last 10 trials: 2.12221825215" and "Average

Destinations Reached for the last 10 trials: 0.9855555555556”. On a whim, I decided to try one other tuning,  $\epsilon=0.2$ ,  $\alpha=0.2$ , and  $\gamma=0.2$ , which resulted in “Average Rewards per Turn for the last 10 trials: 2.03261044849” and “Average Destinations Reached for the last 10 trials: 0.9955555555556”. So I ended up selecting  $\epsilon=0.2$ ,  $\alpha=0.2$ , and  $\gamma=0.2$  as my final model. The graph below shows a summary of those 100 runs.




Figure 1: Summary Plot for  $\epsilon=0.2$ ,  $\alpha=0.2$ , and  $\gamma=0.2$

My agent is currently getting close to a perfect 1.0 for its average destinations reached, but browsing through the logs shows that it does occasionally take illegal actions in its final 10 trials. From my most recent run there were two moves that incurred negative rewards:

Trial 91, inputs = {'light': 'red', 'oncoming': None, 'right': None, 'left': 'forward'}, action = forward, reward = -1.0

Trial 96, inputs = {'light': 'green', 'oncoming': None, 'right': 'forward', 'left': None}, action = forward, reward = -0.5

My thinking is that these states are either encountered too rarely (but considering my now reduced state space, I think this is unlikely. My estimation is that the average run has 1300 moves, which should be enough for a state space of 48.) or the forecasting component of the Q equation is causing it. So I’m expecting that the car is learning to take illegal moves anytime a reached destination inflates

a specific state outside of its otherwise correct policy. I think the solution to this would be to introduce a decay on the alpha, same as what I've already now implemented on the epsilon.

An optimal policy for this problem would have a perfect destination count of 1 , and it would never incur any negative rewards. My learner isn't quite there yet, but it's pretty close.